南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

2017-2018 Academic Year
Spring Semester
Final Exam Paper

Course Name：CS307 Dept.：Computer Science Exam Duration： 120 min

Question No.	1	2	3	4	5			
Score								

This exam paper contains 5 independent questions and the score is 100 in total. (Please hand in your exam paper, answer sheet, and your scrap paper to the proctor when the exam ends.)

All printed and manuscript documents allowed, calculators allowed, other electronic devices forbidden.

Don't get stuck on a question – try to organize your time on the basis one point = one minute.

You must write your answers on the exam paper.

You are reminded of academic integrity requirements.
Papers strangely similar will get 0.

Points are indicative. They may be adjusted if one question is failed by far too many people.

Question 1: Quiz (20 points)
Question 2: Spot the mistake (18 points)
Question 3: Sizing (12 points)
Question 4: Query Analysis (25 points)
Question 5: Design (25 points)

Question 1: Quiz (20 points)

Please divide your answers on the exam paper by blocks of five, like this (w, x, y and z represent possible answers):
1-5: x,y,z,w,x
6-10: y,z,w,x,y
and so forth.

1. Suppose that in a transaction an update fires a trigger that inserts a row into a log table. By default, if the update is rolled back, the row in the log table disappears
    a. True
    b. False
2. A view cannot return duplicate rows
    a. True
    b. False
3. A query on a view can never use the indexes on the table(s) on which the view is based
    a. True
    b. False
4. You can create triggers on views
    a. True
    b. False
5. A logical backup always requires the database to be shut down
    a. True
    b. False
6. A physical backup always requires the database to be shut down
    a. True
    b. False
7. True disaster recovery implies data replication between at least two sites
    a. True
    b. False
8. In a read consistent model
    a. If a session, in a transaction, updates a table and then queries the row that it changed, the select returns the value before the transaction started.
    b. The session that started the transactions sees the changes, but

other sessions don't see them before they are committed.

c. The session that started the transactions sees the changes, but there is a lock and other sessions wait until the transaction is committed if they query a value that was changed.
d. All sessions see the new values.

9. (2 points)If there is no primary key nor unique constraint, a statement such as:

```
insert into T select * from T
```

a. Fails because it's not consistent
b. Doubles the number of rows in table T
c. Inserts rows into T until one runs out of space inside the database

10. (2 points) Suppose that the right to select from a table T was granted to role R, and that the role R was granted to a user U. After a database administrator (or the owner of T) runs

```
revoke select on T from U
```

a. User U can no longer query from T
b. User U can still query from T

11. With most databases systems, an exclusive lock on a row prevents other sessions from reading the row.

a. True
b. False

12. A SELECT statement generates no locks

a. True
b. False

13. (2 points) If I want to remove a film with movieid 123456 from the movie database (reminder: tables are countries, movies that references countries, alt_titles that references movies, people, credits that references movies and people, there is no `ON DELETE CASCADE` on any foreign key), the correct sequence of operations is:

a. delete from movies where movieid = 123456;
   delete from alt_titles where movieid = 123456;
   delete from credits where movieid = 123456;
b. delete from countries
   where country_code = (select country
   from movies
   where movieid = 123456);
   delete from credits where movieid = 123456;
   delete from alt_titles where movieid = 123456;
   delete from movies where movieid = 123456;
c. delete from people
   where peopleid in (select peopleid

from credits
where movieid = 123456);
delete from credits where movieid = 123456;
delete from alt_titles where movieid = 123456;
delete from movies where movieid = 123456;

d. delete from credits where movieid = 123456;
delete from alt_titles where movieid = 123456;
delete from movies where movieid = 123456;

14. (3 points)You have the following table (no primary key):

```
select * from test;
VAL
----
   2
   2
   2
   1
   1
```

The query "`select sum(2) from test`" returns:

a. null
b. 0
c. 3
d. 6
e. 10
f. An error

15. A condition WHERE COL LIKE ... can never use an index on COL.

a. True
b. False

Question 2: Spot the mistake (18 points)

2.1 Don't believe everything you find on the Web (8 points)

A number of websites supply "technical interview questions". The following questions and answers come from one such website:

How to find a duplicate record?

1. *duplicate records with one field*
2. *duplicate records with more than one field*

Answer.

1. *duplicate records with one field*

```
SELECT name, COUNT(email)
FROM users
GROUP BY email
HAVING COUNT(email) > 1
```

2. *duplicate records with more than one field*

```
SELECT name, email, COUNT(*)
FROM users
GROUP BY name, email
HAVING COUNT(*) > 1
```

What is wrong in the answer?

2.2 Desperate Project Manager (10 points)

The following query was posted on a forum by a project manager (I have slightly modified it, but very little) and was posted with the following comment "As my developer doesn't quite master SQL and its subtleties, I'm posting this query here with the hope of getting some help. This query returns the 10 videos having the greatest number of categories in common with the video being watched (in this example video #81). This query takes 5 seconds with 2700 videos on a powerful server, we find it rather slow."

```
SELECT DISTINCT
       video_id,
       video_type,
       video_title,
       video_description,
       video_idPartner,
       video_urlMini,
       video_dateValid,
       partner_valid,
       partner_redirection,
       ( SELECT COUNT(Y.v_belongs_c_idVideo) AS NbSimilar
          FROM  v_belongs_c  Y
          WHERE Y.v_belongs_c_idVideo=81
            AND Y.v_belongs_c_idCategory IN
              (SELECT Z.v_belongs_c_idCategory
               FROM v_belongs_c Z
               WHERE Z.v_belongs_c_idVideo=video_id)) as Counter,
        ( SELECT category_singular
          FROM category,
               v_belongs_c X
```

```
            WHERE X.v_belongs_c_idVideo=video_id
              AND X.v_belongs_c_default=1
              AND category_id=X.v_belongs_c_idCategory )
                      as category_singular
FROM category,
     v_belongs_c A,
     video
     LEFT JOIN partner
            ON video_idPartner=partner_id
WHERE (A.v_belongs_c_idCategory IN
          (SELECT W.v_belongs_c_idCategory
           FROM v_belongs_c W
           WHERE W.v_belongs_c_idVideo=81)
       AND video_id=A.v_belongs_c_idVideo)
 AND (video_idPartner=0
      OR (partner_valid=1
          AND partner_redirection<>1))
 AND video_valid=1
 AND video_id <> 81
ORDER BY Counter DESC
LIMIT 10
```

WITHOUT TRYING TO REWRITE THE QUERY, can you point to a major issue in the query, and how the developer “fixed” it (hint: badly).

Question 3: Sizing (12 points)

This company wants to store in a specialized datamart information that comes from telephony systems; telephony systems are mostly today computers routing calls, which include an internal database, and generate on demand “Call Detail Records” which contain information about every phone call passed through the system (number of the caller, number called, number that answered, start time, duration, bytes of data transmitted, etc.) As most calls are passed over IP, the goal is overall to check how much bandwidth is used at different times of day, monitor the use of gateways and measure the impact of some hardware failure (would half the company be left in the dark?), predict the impact of a video-conferencing system ... and get some quantitative productivity measurements for some departments where the phone is the main tool, such as customer support and telemarketing.

Call Detail Records, or CDRs, contain on average 200 bytes. The telephony system in this (big) multinational company serves a wide geographic area and there are overall around 200,000 calls a day. You are reminded that data blocks

contain a header and various overhead, and that raw data will use around 15% more storage when inserted into a database table. Additionally, an estimated 60% of table storage will be used by indexes. This will be the "big table" in the database, the remainder (data dictionary, other tables, system storage) can be estimated at around 400M.
This system is a decision support system and will not be used 24x7. Additionally, data is obtained from the telephony system computers (Cisco call them "Call Managers") that can generate on demand CDR files for the past 10 days. As a consequence, a daily cold backup every night is considered to be quite sufficient, as it takes about 20 minutes to load and process daily data from a CDR file. However, many batch processes are running during the night and the maintenance window is narrow.
We have relatively slow disks that allow transferring about 60M/s. We want to be able to backup or restore the database in 30 minutes or less. How long can we keep data online before archiving it? In other words, how far back in time will we be able to go if we want to keep the size such as we can backup or restore the database in under 30 minutes?

Question 4: Query analysis (25 points)

The following figure describes a database used to store information about the states and union territories of India, as well as the official languages in use.

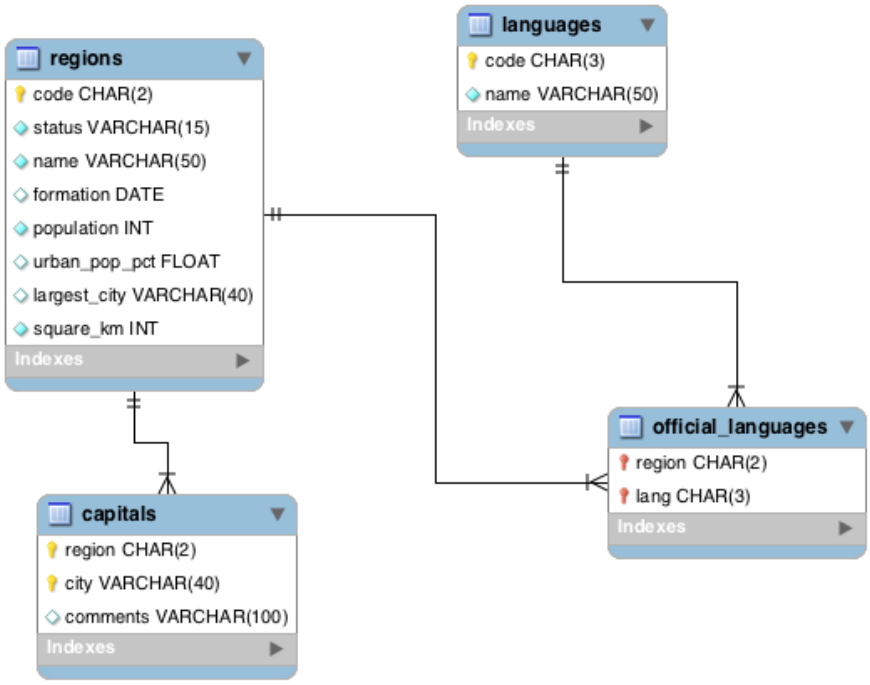

Table Descriptions

regions

code	Two-letter code	Official code (in uppercase) for the state or the region. Always populated.
status	String	Tells whether the region is a state or a union territory, contains either *State* or *Union Territory*. Always populated.
name	String	Name of the region. Always populated.
formation	Date	Date of formation for states. Empty for union territories.
population	Integer	Population of the region. Always populated.
urban_pop_pct	Integer	Percentage of the population of the region that lives in cities.
largest_city	String	Name of the largest city in the region. Empty if the largest city is the capital.
square_km	Integer	Area of the region, expressed in square kilometers. Always populated.

capitals

region	Two-letter code	Official code in uppercase for the region, matches the code column in the region table. Always populated.
city	String	Name of the capital. Always populated.
comments	String	Special comment for unusual cases, such as capitals shared by two regions. Usually empty.

official_languages

region	Two-letter code	Official code in uppercase for the region, matches the code column in the region table. Always populated.
lang	Three-letter code	Lowercase code of a language, matches the code column in the languages table. Always populated.

languages

code	Three-letter code	Standard code in lowercase for a language. Always populated.
name	String	Name of the language. Always populated.

For the following questions, give the identifier of the queries that correctly answer the questions, "None" if no query gives the correct answer. The queries were written with SQLite and are assumed to run.

A	What is the total population of states and territories where Hindi is an official language?

A.1	`SELECT SUM(r.population)`
	`FROM regions r`
	`  INNER JOIN official_languages ol`
	`  ON ol.region = r.code`
	`  INNER JOIN languages l`
	`  ON l.code = ol.lang`
	`WHERE l.name = 'Hindi'`

A.2	`select la.name, reg.population`
	`from (select re.code, re.population`
	`  from regions re`
	`  group by re.code) reg`
	`  join official_languages of`
	`  on of.region = reg.code`
	`  join languages la`
	`  on la.code = of.lang`
	`where la.name = 'Hindi'`

A.3	`select sum(population)`
	`from regions r`
	`  join official_languages ol`
	`  on r.code = ol.region`
	`  join languages l`
	`  on ol.lang = l.code`
	`where l.name = 'Hindi'`

B	What are the capitals that aren't the only capital of a single state or territory (some capitals are shared, some states have several capitals) - the query should return TWO columns 1) Comma-separated list of capitals (when there are several capitals 2) Comma-separated list of states (when several states)

B.1	`SELECT group_concat(c.city),group_concat(c.region)`
	`FROM capitals c`
	`GROUP BY c.city, c.region`
	`HAVING count(c.city) > 1`
	`  OR count(c.region) > 1`

B.2	`select group_concat(ca.city) cities,`
	`  group_concat(reg.name) regions`
	`from (select re.code, re.name`
	`  from regions re`
	`  group by re.code) reg`
	`  join capitals ca`
	`  on ca.region = reg.code`
	`where ca.comments != ''`

B.3	`select r.name, c.city, count(*) c_count`
	`from regions r`
	`  join capitals c`
	`  on r.code = c.region`
	`group by c.city`
	`having c_count > 1`
	`union`
	`select r.name, c.city, count(*) s_count`
	`from regions r`
	`  join capitals c`
	`  on r.code = c.region`
	`group by r.name`
	`having s_count > 1`

C	What are the name, population, and number of languages of the state(s) or region(s) with the most official languages?

C.1	`SELECT r.name, SUM(population), COUNT(ol.lang) AS langCount`
	`FROM regions r`
	`  INNER JOIN official_languages ol`
	`  ON ol.region = r.code`
	`WHERE langCount = (SELECT COUNT(ol.lang) AS langCount2`
	`  FROM official_languages ol`
	`  GROUP BY ol.region`
	`  ORDER BY langCount2 desc`
	`  LIMIT 1)`
	`GROUP BY r.name`
	`ORDER BY langCount desc`

C.2	`select re.name, re.population, off.numOfLang`
	`from (select of.region, count(lang) numOfLang`
	`  from official_languages of`
	`  group by of.region) off`
	`  join regions re`
	`  on re.code = off.region`

C.3	`select r.name, r.population, count(*) langs`
	`from regions r`
	`  join official_languages ol`
	`  on r.code = ol.region`
	`  left join languages l`
	`  on ol.lang = l.code`
	`group by r.code`
	`having langs = (select max(n_langs)`
	`  from (select count(*) n_langs`
	`  from regions r`
	`  join official_languages ol`
	`  on r.code = ol.region`
	`  group by r.code))`

D	What are the States (<u>not</u> Union Territories) where English is NOT an official language? Display their names in alphabetical order.

D.1	`SELECT DISTINCT r.name`
	` FROM regions r`
	` INNER JOIN official_languages ol`
	`  ON ol.region = r.code`
	` INNER JOIN languages l`
	`  ON l.code = ol.lang`
	` WHERE r.status = 'State'`
	`  AND l.name <> 'English'`
	` ORDER BY r.name`

D.2	`select re.name`
	`from (select of.region, of.lang`
	`from official_languages of`
	`group by of.region) off`
	`join regions re`
	`on re.code = off.region`
	`join languages la`
	`on la.code = off.lang`
	`where re.status = 'State'`
	`  and la.name != 'English'`

(D.3 follows on next page)

D.3	`select r.name, l.name`
	`from regions r`
	`  join official_languages ol`
	`  on r.code = ol.region`
	`  join languages l`
	`  on ol.lang = l.code`
	`group by r.name`
	`having l.name != 'English'`
	`  and r.status = 'State'`
	`order by r.name asc`

Question 5: Design (25 points)

We want to create a music database. Usually you have a relationship between music album and artists. We'll just focus here on artists. The problem of artists is that an album can be credited to an individual artist (say " Jacky Cheung (张学友)" or a band (say "Phoenix Legend (凤凰传奇)" or "Beyond"); so a band can also be considered an artist. We may want for individuals to record date of birth and (possibly) of death, and for bands the date of creation and (possibly) the date when the band was officially disbanded. Additionally, we may want to record who are the members of a pop band (perhaps not of a symphonic orchestra), and band members may change over time. Band members can also move to other bands, or start a solo career.

5.1. (20 points) How would you model what is described above?

5.2. (5 points) Would you make the artist identifier an attribute of an album, or would you use a relationship table (artistid, albumid)? Justify.